

Chapter 3 summary

- An operating system executes a variety of programs:
 - Batch system – **jobs**
 - Time-shared systems – **user programs** or **tasks**
- Program is **passive** entity stored on disk (*executable file*); process is **active**.
 - Program becomes process when executable file loaded into memory.
- One program can be **several** processes.
- Process includes multiple parts:
 - The program code, also called **text section**.
 - Current activity involving **program counter**, processor registers.
 - **Stack** containing temporary data.
 - **Data section** containing global variables.
 - **Heap** containing memory dynamically allocated during run time.
- As a process executes, it changes **state**:
 - **new**: The process is being created.
 - **running**: Instructions are being executed.
 - **waiting**: The process is waiting for some event to occur.
 - **ready**: The process is waiting to be assigned to a processor.
 - **terminated**: The process has finished execution.
- Process Control Block (PCB):
 - **Process state** – running, waiting, etc.
 - **Program counter** – location of instruction to next execute.
 - **CPU registers** – contents of all process-centric registers.
 - **CPU scheduling information** – priorities, scheduling queue pointers.
 - **Memory-management information** – memory allocated to the process.
 - **Accounting information** – CPU used, clock time elapsed since start, time limits.
 - **I/O status information** – I/O devices allocated to process, list of open files.
- Process Scheduling:
 - Maximize CPU use, quickly switch processes onto CPU for time sharing.
 - **Process scheduler** selects among available processes for next execution on CPU.
 - Maintains **scheduling queues** of processes:
 - **Job queue** – set of all processes in the system.
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute.
 - **Device queues** – set of processes waiting for an I/O device.
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue. (may be slow)

- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU. (must be fast)
- The long-term scheduler controls the **degree of multiprogramming**.
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O, many short CPU bursts.
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts.
- Long-term scheduler strives for good **process mix**.
- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease.
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**.
- Context switch:
 - When CPU switches to another process, the system must **save the state** of the old process and **load the saved state** for the new process via a **context switch**. (PCB)
 - Context-switch time is **overhead**; the system does no useful work while switching.
 - The more complex the OS and the PCB, longer the **context switch**.
 - Time dependent on **hardware support**.
- **Parent** process create **children** processes, which, in turn create other processes, forming a tree of processes.
 - Generally, process identified and managed via a **process identifier (pid)**.
- Resource sharing options:
 - Parent and children share all resources.
 - Children share subset of parent's resources.
 - Parent and child share no resources.
- Execution options:
 - Parent and children execute concurrently.
 - Parent waits until all or some children terminate.
- Address space:
 - Child duplicate of parent program.
 - Child has a program loaded into it.
- UNIX examples:
 - **fork()** system call creates new process.
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program.

- Process executes last statement and asks the operating system to delete it **exit()**.
- Parent may terminate execution of its children processes **abort()** for a variety of reasons:
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - If parent is exiting:
 - Some operating systems do not allow child to continue if its parent terminates.
 - All children must also be terminated - **cascading termination**.
- Processes within a system may be **independent** or **cooperating**.
- Independent process cannot affect or be affected by other processes.
- Cooperating process can affect or be affected by other processes, including **sharing data**. Also, require an **interprocess communication (IPC)** mechanism.
 - Two models of IPC:
 - Message passing.
 - Shared memory.
- **Producer-Consumer Problem**: producer process produces information that is consumed by a consumer process.
- One solution to allow them to run concurrently, uses a buffer of items (**shared memory**) that can be filled by producer and emptied by consumer:
 - **Unbounded-buffer** places no practical limit on the size of the buffer.
 - **Bounded-buffer** assumes that there is a fixed buffer size.
- Message passing – processes communicate with each other without resorting to shared variables.
 - IPC facility provides at least two operations:
 - **send(message)** – message size fixed or variable.
 - **receive(message)**.
 - If P and Q wish to communicate, they need to:
 - establish a **communication link** between them.
 - exchange messages via send/receive.
- Direct communication:
 - Properties of communication link:
 - Links are established **automatically**.
 - A link is associated with exactly **one** pair of communicating processes.
 - Between each pair there exists exactly **one** link.
 - The link may be **unidirectional** but is usually **bi-directional**.

- Indirect communications:
 - Messages are directed and received from **mailboxes** (also referred to as ports)
 - Each mailbox has a unique id.
 - Processes can communicate only if they share a mailbox.
 - Properties of communication link:
 - Link established only if processes share a **common mailbox**.
 - A link may be associated with **many** processes.
 - Each pair of processes may share **several** communication links.
 - Link may be **unidirectional** or **bi-directional**.
- Synchronization:
 - Message passing may be either blocking or non-blocking.
 - Blocking is considered **synchronous**.
 - **Blocking send** has the sender block until the message is received.
 - **Blocking receive** has the receiver block until a message is available.
 - Non-blocking is considered **asynchronous**.
 - **Non-blocking send** has the sender send the message and continue.
 - **Non-blocking receive** has the receiver receive a valid message or null.
 - If both send and receive are blocking, we have a **rendezvous**.
- Buffering:
 - Queue of messages attached to the link; implemented in one of three ways:
 - **Zero capacity – 0 messages**
 - Sender must wait for receiver (rendezvous).
 - **Bounded capacity – finite length of n messages**
 - Sender must wait if link full.
 - **Unbounded capacity – infinite length**
 - Sender never waits.
- Shared memory and message passing strategies can be used for communication in client–server systems as well.
- A socket is an endpoint for communication. A socket is identified by an **IP address** concatenated with a **port number**.
- The server waits for incoming client requests by listening to a specified port
 - All ports below 1024 are **well known**, used for standard services
 - To allow a client and server on the same host to communicate, a special IP address 127.0.0.1 (**loopback**) is used to refer to itself.
- Three different types of sockets in java:
 - **Connection-oriented (TCP)**
 - **Connectionless (UDP)**
 - **MulticastSocket class** – data can be sent to multiple recipients.

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
 - The client-side stub locates the server and **marshalls** the parameters.
 - The server-side stub receives this message, unpacks the **marshalled** parameters, and performs the procedure on the server.
- **Pipes**: acts as a channel allowing two processes to communicate.
- **Ordinary Pipes** allow communication in standard producer-consumer style.
 - Allow only **unidirectional** communication.
- Require **parent-child** relationship between communicating processes.
- UNIX treats it as a special type of file and can be accessed using **read()** and **write()** system calls.
- Named pipes are more powerful tool than ordinary pipes:
 - Communication can be bidirectional.
 - No parent–child relationship is required.
 - Several processes can use it for communication.
- Both UNIX and Windows systems support it, although implementation details differ greatly.