

## Chapter 9 summary

- Code needs to be in memory to execute, but entire program rarely used.
- Entire program code not needed at same time.
- Consider ability to execute partially loaded program:
  - Program no longer constrained by limits of physical memory.
  - Each program takes less memory while running → more programs run at the same time → Increased CPU utilization and throughput.
  - Less I/O needed to load or swap programs into memory → each user program runs faster.
- **Virtual memory** – separation of user logical memory from physical memory.
  - Only part of the program needs to be in memory for execution.
  - Logical address space can therefore be **much larger** than physical address space.
  - Allows address spaces to be shared by several processes.
  - Allows for more efficient process creation.
  - More programs running concurrently.
  - Less I/O needed to load or swap processes.
- **Virtual address space** – logical view of how process is stored in memory.
- **Demand Paging:**
  - Could bring entire process into memory at load time or bring a page into memory only when it is needed.
    - Less I/O needed, no unnecessary I/O.
    - Less memory needed.
    - Faster response.
    - More users.
  - Like **paging system** with swapping.
  - Page is needed – reference to it:
    - invalid reference (accessing page that is not related to the current program) → abort.
    - not-in-memory → bring to memory.
  - **Lazy swapper** – never swaps a page into memory unless page will be needed.
    - Swapper that deals with pages is a pager.
  - Pages that are always needed are **memory resident**.
  - If page needed and not memory resident:
    - Need to detect and load the page into memory from storage.
- **Valid-Invalid Bit**
  - **v** → in-memory – **memory resident**
  - **i** → not-in-memory

- During MMU address translation, if valid–invalid bit in page table entry is i → **page fault**.
- **Page Fault (steps in handling page fault):**
  - Operating system looks at the process's internal table to decide:
    - Invalid reference → abort
    - Just not in memory
  - Find free frame.
  - Swap page into frame via scheduled disk operation.
  - Reset tables to indicate page now in memory (Set validation bit = v ).
  - Restart the instruction that caused the page fault.
- **Aspects of demand paging:**
  - **Pure demand paging** – start process with no pages in memory.
  - **Locality of reference** – Some localities that their references are known and are used continuously.
  - Hardware support needed for demand paging:
    - Page table with valid / invalid bit.
    - Secondary memory (swap device with swap space).
    - Ability to restart any instruction after a page fault.
- **Instruction restart:**
  - Must be able to restart the instruction in **exactly** in the same place and state.
  - Consider an instruction that could access several different locations.
    - Take care if there is an overlap!
- **Performance of demand paging:**
  - Service the interrupt – careful coding means just several hundred instructions needed.
  - Read the page – lots of time.
  - Restart the process – again just a small amount of time.
  - Page Fault Rate:  $0 \leq p \leq 1$
  - Effective Access Time (EAT):
    - $EAT = (1 - p) \times \text{memory access} + p (\text{page fault overhead} + \text{swap page out} + \text{swap page in})$
    - Page fault ↑ memory access time ↑ performance ↓
- **Demand paging optimization:**
  - Swap space I/O faster than file system I/O even if on the same device.
  - Demand pages for program binary files – they are never modified.
- **Copy-on-write:**
  - Allow both parent and child processes to initially **share** the same pages in memory.

- if either process writes to a shared page, a copy of the shared page is created.
  - In general, free frames are allocated from a pool of zero-fill-on-demand pages.
- **Page replacement** – find some page in memory, but not really in use, page it out.
- Same page may be brought into memory several times.
- Page replacement:
  - Prevent over-allocation of memory by modifying page-fault service routine to include page replacement.
  - Use modify (dirty) bit to reduce overhead of page transfers.
    - Only modified pages are written to disk.
  - Page replacement completes separation between logical memory and physical memory.
- Basic page replacement:
  - Find the location of the desired page on disk.
  - Find a free frame:
    - If there is a free frame, use it.
    - If there is no free frame, use a page replacement algorithm to select a victim frame.
    - Write victim frame to disk if dirty.
  - Bring the desired page into the (newly) free frame; update the page and frame tables.
  - Continue the process by restarting the instruction that caused the trap.

(Note: if no frames are free, two-page transfers are required for page fault – increasing EAT (Effective Access Time))
- Frame-allocation algorithm determines how many frames to give each process.
  - Frames ↑ page fault ↓
- Evaluate each algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string.
  - String is just page numbers, not full addresses.
  - Repeated access to the same page does not cause a page fault.
- **First-In-First-Out (FIFO) Algorithm:** *\*Important example in slide 34\**
  - Belady's Anomaly: Adding more frames can cause more page faults!
- **Optimal algorithm:** *\*Important example in slide 37\**
  - Replace page that will not be used for longest period of time.
  - Used for measuring how well your algorithm performs.
- **Least Recently Used (LRU) Algorithm:** *\*Important example in slide 38\**
  - Use past knowledge rather than future.
  - Replace page that has not been used in the most amount of time.

- Better than FIFO but worse than OPT.
- Counter implementation:
  - Every page-table entry has a time-of-use field; whenever a reference to a page is made, the contents of the clock are copied to this field its entry.
- Stack implementation:
  - Keep a stack of page numbers, whenever a page is referenced, it is removed and put on top of the stack.
  - Least recently used page is at bottom of the stack.
- **LRU Approximation Algorithms:**
  - Reference bit:
    - With each page associate a bit, initially = 0
    - When page is referenced bit set to 1
    - Replace any with reference bit = 0 (if one exists)
      - However, we do not know the order.
  - Second-chance algorithm:
    - Generally, FIFO plus hardware-provided reference bit
    - Clock replacement
    - If page to be replaced has:
      - Reference bit = 0 => replace it
      - Reference bit = 1 then:
        - set reference bit 0, leave page in memory.
        - replace next page, subject to same rules.
- **Enhanced Second-Chance Algorithm:**
  - Improve algorithm by using reference bit and modify bit (if available) in concert.
  - Take ordered pair (reference, modify)
    - (0, 0) neither recently used not modified – best page to replace.
    - (0, 1) not recently used but modified – not quite as good, must write out before replacement.
    - (1, 0) recently used but clean – probably will be used again soon.
    - (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement.
- **Counting Algorithms:**
  - Keep a counter of the number of references that have been made to each page.
  - Least Frequently Used (LFU) Algorithm: replaces the page with smallest count
    - A problem when a page is used heavily during initial phase of a process but then is never used again.
  - Most Frequently Used (MFU) Algorithm: replaces the page with largest count
- **Page-Buffering Algorithms:**
  - Keep a pool of free frames, always and keep a list of modified pages.

- **Allocation of frames:**

- Each process needs minimum number of frames.
- Maximum of course is total frames in the system.
- Fixed Allocation:
  - Fixed equal allocation:
    - Equal number of frames to each process.
  - Fixed proportional allocation:
    - Allocate according to the size of process.
    - Dynamic as degree of multiprogramming, process sizes change.
- Priority Allocation:
  - Use a proportional scheme wherein ratio of frames depends on priorities of processes rather than on relative sizes of processes or on a combination of size and priority.

- **Global vs. Local Allocation:**

- Global replacement – a process selects a replacement frame from the set of all frames; one process can take a frame from another.
- Local replacement – each process selects from only its own set of allocated frames.
  - More consistent per-process performance.
  - But possibly underutilized memory.

- **Thrashing:**

- If a process does not have “enough” pages, the page-fault rate is very high.
  - Page fault to get page.
  - Replace existing frame.
  - But quickly need replaced frame back.
  - This leads to:
    - Low CPU utilization.
    - Operating system thinking that it needs to increase the degree of multiprogramming.
    - Another process added to the system.
- Thrashing → a process is busy swapping pages in and out.